

Support Vector Machine Classifiers

Outline

- Support Vector Machines for Classification
 - Linear Discrimination
 - Nonlinear Discrimination
- SVM Mathematically
- Extensions
- Data Classification
- Kernel Functions


Definition

- ‘Support Vector Machine is a system for efficiently training **linear learning machines** in **kernel-induced feature spaces**, while respecting the insights of **generalisation** theory and exploiting **optimisation** theory.’


- **AN INTRODUCTION TO SUPPORT VECTOR MACHINES** (and other kernel-based learning methods)

N. Cristianini and J. Shawe-Taylor
Cambridge University Press
2000 ISBN: 0 521 78019 5

- Kernel Methods for Pattern Analysis
John Shawe-Taylor & Nello Cristianini
Cambridge University Press, 2004



The Scalar Product



$$\mathbf{a} \cdot \mathbf{b} = |\mathbf{a}| |\mathbf{b}| \cos \theta$$

The scalar or dot product is, in some sense, a measure of **Similarity**

Decision Function for binary classification

$$f(x) \in \mathbf{R}$$

$$f(x_i) \geq 0 \Rightarrow y_i = 1$$

$$f(x_i) < 0 \Rightarrow y_i = -1$$

Support Vector Machines

- SVMs pick **best** separating hyperplane according to some criterion
 - e.g. maximum margin
- Training process is an **optimisation**
- Training set is effectively reduced to a relatively small number of **support vectors**

Feature Spaces

- We may separate data by mapping to a higher-dimensional feature space
 - The feature space may even have an infinite number of dimensions!
- We need not **explicitly** construct the new feature space

Kernels

- We may use Kernel functions to **implicitly** map to a new feature space
- Kernel fn:

$$K(\mathbf{x}_1, \mathbf{x}_2) \in \mathbf{R}$$

- Kernel must be equivalent to an **inner product** in some feature space

Example Kernels


Linear: $\langle \mathbf{x} \cdot \mathbf{z} \rangle$

Polynomial: $P(\langle \mathbf{x} \cdot \mathbf{z} \rangle)$

Gaussian: $\exp\left(-\|\mathbf{x} - \mathbf{z}\|^2 / \sigma^2\right)$


Perceptron Revisited: Linear Separators

- Binary classification can be viewed as the task of separating classes in feature space:




$$f(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x} + b)$$


Which of the linear separators is optimal?




Best Linear Separator?




Best Linear Separator?




Best Linear Separator?




Best Linear Separator?



Find Closest Points in Convex Hulls




Plane Bisect Closest Points




Classification Margin

- Distance from example data to the separator is $r = \frac{\mathbf{w}^T \mathbf{x} + b}{\|\mathbf{w}\|}$
- Data closest to the hyperplane are *support vectors*.
- *Margin* ρ of the separator is the width of separation between classes.



Maximum Margin Classification


- Maximizing the margin is good according to intuition and theory.
- Implies that only support vectors are important; other training examples are ignorable.



Statistical Learning Theory


- Misclassification error and the function complexity bound generalization error.
- Maximizing margins minimizes complexity.
- “Eliminates” **overfitting**.
- Solution depends only on *Support Vectors* not number of attributes.

Margins and Complexity



Skinny margin
is more flexible
thus more complex.

Margins and Complexity



Fat margin
is less complex.

Linear SVM Mathematically

- Assuming all data is at distance larger than 1 from the hyperplane, the following two constraints follow for a training set $\{(\mathbf{x}_i, y_i)\}$

$$\mathbf{w}^T \mathbf{x}_i + b \geq 1 \quad \text{if } y_i = 1$$

$$\mathbf{w}^T \mathbf{x}_i + b \leq -1 \quad \text{if } y_i = -1$$

- For support vectors, the inequality becomes an equality; then, since each example's distance from the

- hyperplane is $r = \frac{\mathbf{w}^T \mathbf{x} + b}{\|\mathbf{w}\|}$ the margin is: $\rho = \frac{2}{\|\mathbf{w}\|}$

Linear SVMs Mathematically (cont.)

- Then we can formulate the *quadratic optimization problem*:

Find \mathbf{w} and b such that

$$\rho = \frac{2}{\|\mathbf{w}\|} \quad \text{is maximized and for all } \{(\mathbf{x}_i, y_i)\}$$
$$\mathbf{w}^T \mathbf{x}_i + b \geq 1 \text{ if } y_i = 1; \quad \mathbf{w}^T \mathbf{x}_i + b \leq -1 \text{ if } y_i = -1$$

A better formulation:

Find \mathbf{w} and b such that

$$\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} \quad \text{is minimized and for all } \{(\mathbf{x}_i, y_i)\}$$

$$y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$$

Solving the Optimization Problem

Find \mathbf{w} and b such that

$$\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} \text{ is minimized and for all } \{(\mathbf{x}_i, y_i)\}$$
$$y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$$

- Need to optimize a *quadratic* function subject to *linear* constraints.
- Quadratic optimization problems are a well-known class of mathematical programming problems, and many (rather intricate) algorithms exist for solving them.
- The solution involves constructing a *dual problem* where a *Lagrange multiplier* α_i is associated with every constraint in the primary problem:

Find $\alpha_1 \dots \alpha_N$ such that

$$\mathbf{Q}(\boldsymbol{\alpha}) = \sum \alpha_i - \frac{1}{2} \sum \sum \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j \text{ is maximized and}$$

$$(1) \sum \alpha_i y_i = 0$$

$$(2) \alpha_i \geq 0 \text{ for all } \alpha_i$$

The Optimization Problem Solution

- The solution has the form:

$$\mathbf{w} = \sum \alpha_i y_i \mathbf{x}_i \quad b = y_k - \mathbf{w}^T \mathbf{x}_k \text{ for any } \mathbf{x}_k \text{ such that } \alpha_k \neq 0$$


- Each non-zero α_i indicates that corresponding \mathbf{x}_i is a support vector.
- Then the classifying function will have the form:

$$f(\mathbf{x}) = \sum \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b$$

- Notice that it relies on an *inner product* between the test point \mathbf{x} and the support vectors \mathbf{x}_i – we will return to this later!
- Also keep in mind that solving the optimization problem involved computing the inner products $\mathbf{x}_i^T \mathbf{x}_j$ between all training points!

Soft Margin Classification

- What if the training set is not linearly separable?
- *Slack variables* ξ_i can be added to allow misclassification of difficult or noisy examples.



Soft Margin Classification Mathematically

- The old formulation:

Find \mathbf{w} and b such that

$$\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} \text{ is minimized and for all } \{(\mathbf{x}_i, y_i)\}$$
$$y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$$

- The new formulation incorporating slack variables:

Find \mathbf{w} and b such that

$$\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum \xi_i \text{ is minimized and for all } \{(\mathbf{x}_i, y_i)\}$$
$$y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i \quad \text{and} \quad \xi_i \geq 0 \text{ for all } i$$

- Parameter C can be viewed as a way to control overfitting.

Soft Margin Classification – Solution

- The **dual problem** for soft margin classification:

Find $\alpha_1 \dots \alpha_N$ such that

$Q(\boldsymbol{\alpha}) = \sum \alpha_i - \frac{1}{2} \sum \sum \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$ is maximized and

(1) $\sum \alpha_i y_i = 0$

(2) $0 \leq \alpha_i \leq C$ for all α_i

- Neither slack variables ξ_i nor their Lagrange multipliers appear in the dual problem!
- Again, \mathbf{x}_i with non-zero α_i will be **support vectors**.
- Solution to the dual problem is:

$$\mathbf{w} = \sum \alpha_i y_i \mathbf{x}_i$$

$$b = y_k (1 - \xi_k) - \mathbf{w}^T \mathbf{x}_k \text{ where } k = \underset{k}{\operatorname{argmax}} \alpha_k$$

But neither \mathbf{w} nor b are needed explicitly for classification!

$$f(\mathbf{x}) = \sum \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b$$

Theoretical Justification for Maximum Margins

- Vapnik has proved the following:

The class of optimal linear separators has VC dimension h bounded from above as

$$h \leq \min \left\{ \left\lceil \frac{D^2}{\rho^2} \right\rceil, m_0 \right\} + 1$$

where ρ is the margin, D is the diameter of the smallest sphere that can enclose all of the training examples, and m_0 is the dimensionality.

- Intuitively, this implies that regardless of dimensionality m_0 we can minimize the VC dimension by maximizing the margin ρ .
- Thus, complexity of the classifier is kept small regardless of dimensionality.

Linear SVMs: Overview

- The classifier is a *separating hyperplane*.
- Most “important” training points are support vectors; they define the hyperplane.
- Quadratic optimization algorithms can identify which training points \mathbf{x}_i are support vectors with **non-zero Lagrangian multipliers α_i** .
- Both in the dual formulation of the problem and in the solution training points appear only inside inner products:

Find $\alpha_1 \dots \alpha_N$ such that

$\mathbf{Q}(\boldsymbol{\alpha}) = \sum \alpha_i - \frac{1}{2} \sum \sum \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$ is maximized and


(1) $\sum \alpha_i y_i = 0$

(2) $0 \leq \alpha_i \leq C$ for all α_i


$$f(\mathbf{x}) = \sum \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b$$

Non-linear SVMs


- Datasets that are linearly separable with some noise work out great:



- But what are we going to do if the dataset is just too hard?



- How about... mapping data to a higher-dimensional space:



Nonlinear Classification


$$x = [a, b]$$

$$x \square w = w_1 a + w_2 b$$

↓


$$\theta(x) = [a, b, ab, a^2, b^2]$$

$$\theta(x) \square w = w_1 a + w_2 b + w_3 ab + w_4 a^2 + w_5 b^2$$



Non-linear SVMs: Feature spaces

- General idea: the original feature space can always be mapped to some higher-dimensional feature space where the training set is separable:



The “Kernel Trick”

- The linear classifier relies on inner product between vectors $K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$
- If every datapoint is mapped into high-dimensional space via some transformation $\Phi: \mathbf{x} \rightarrow \phi(\mathbf{x})$, the inner product becomes:

$$K(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$$

- A *kernel function* is some function that corresponds to an inner product into some feature space.
- Example:

2-dimensional vectors $\mathbf{x} = [x_1 \ x_2]$; let $K(\mathbf{x}_i, \mathbf{x}_j) = (1 + \mathbf{x}_i^T \mathbf{x}_j)^2$,

Need to show that $K(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$:

$$\begin{aligned} K(\mathbf{x}_i, \mathbf{x}_j) &= (1 + \mathbf{x}_i^T \mathbf{x}_j)^2 = 1 + x_{i1}^2 x_{j1}^2 + 2 x_{i1} x_{j1} x_{i2} x_{j2} + x_{i2}^2 x_{j2}^2 + 2 x_{i1} x_{j1} + 2 x_{i2} x_{j2} = \\ &= [1 \ x_{i1}^2 \ \sqrt{2} x_{i1} x_{i2} \ x_{i2}^2 \ \sqrt{2} x_{i1} \ \sqrt{2} x_{i2}]^T [1 \ x_{j1}^2 \ \sqrt{2} x_{j1} x_{j2} \ x_{j2}^2 \ \sqrt{2} x_{j1} \ \sqrt{2} x_{j2}] = \\ &= \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j), \quad \text{where } \phi(\mathbf{x}) = [1 \ x_1^2 \ \sqrt{2} x_1 x_2 \ x_2^2 \ \sqrt{2} x_1 \ \sqrt{2} x_2] \end{aligned}$$

Advances in Kernel Methods

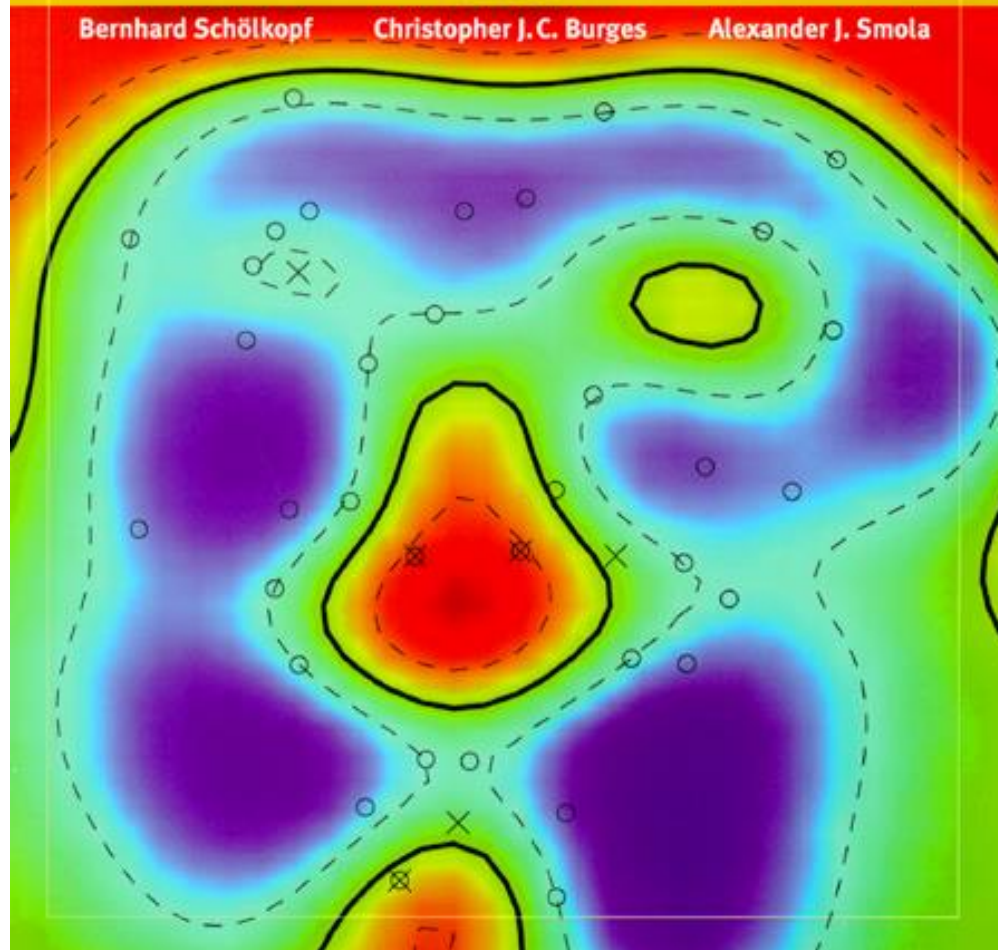
Support Vector Learning

edited by

Bernhard Schölkopf

Christopher J. C. Burges

Alexander J. Smola



Positive Definite Matrices

A square matrix A is *positive definite* if $x^T A x > 0$ for all nonzero column vectors x .

It is *negative definite* if $x^T A x < 0$ for all nonzero x .

It is *positive semi-definite* if $x^T A x \geq 0$.

And *negative semi-definite* if $x^T A x \leq 0$ for all x .

What Functions are Kernels?

- For some functions $K(\mathbf{x}_i, \mathbf{x}_j)$ checking that $K(\mathbf{x}_i, \mathbf{x}_j) = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_j)$ can be cumbersome.
- Mercer's theorem:

Every semi-positive definite symmetric function is a kernel

- Semi-positive definite symmetric functions correspond to a semi-positive definite symmetric Gram matrix:

$K =$

$K(\mathbf{x}_1, \mathbf{x}_1)$	$K(\mathbf{x}_1, \mathbf{x}_2)$	$K(\mathbf{x}_1, \mathbf{x}_3)$...	$K(\mathbf{x}_1, \mathbf{x}_N)$
$K(\mathbf{x}_2, \mathbf{x}_1)$	$K(\mathbf{x}_2, \mathbf{x}_2)$	$K(\mathbf{x}_2, \mathbf{x}_3)$		$K(\mathbf{x}_2, \mathbf{x}_N)$
...
$K(\mathbf{x}_N, \mathbf{x}_1)$	$K(\mathbf{x}_N, \mathbf{x}_2)$	$K(\mathbf{x}_N, \mathbf{x}_3)$...	$K(\mathbf{x}_N, \mathbf{x}_N)$

Examples of Kernel Functions

- Linear: $K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$
- Polynomial of power p : $K(\mathbf{x}_i, \mathbf{x}_j) = (1 + \mathbf{x}_i^T \mathbf{x}_j)^p$
- Gaussian (radial-basis function network): $K(\mathbf{x}_i, \mathbf{x}_j) = e^{-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}}$
- Two-layer perceptron: $K(\mathbf{x}_i, \mathbf{x}_j) = \tanh(\beta_0 \mathbf{x}_i^T \mathbf{x}_j + \beta_1)$

Non-linear SVMs Mathematically

- Dual problem formulation:

Find $\alpha_1 \dots \alpha_N$ such that

$Q(\alpha) = \sum \alpha_i - \frac{1}{2} \sum \sum \alpha_i \alpha_j y_i y_j K(\mathbf{x}_i, \mathbf{x}_j)$ is maximized and

(1) $\sum \alpha_i y_i = 0$

(2) $\alpha_i \geq 0$ for all α_i

- The solution is:

$$f(\mathbf{x}) = \sum \alpha_i y_i K(\mathbf{x}_i, \mathbf{x}) + b$$

- Optimization techniques for finding α_i 's remain the same!

SVM applications

- SVMs were originally proposed by Boser, Guyon and Vapnik in 1992 and gained increasing popularity in late 1990s.
- SVMs are currently among the best performers for a number of classification tasks ranging from text to genomic data.
- SVM techniques have been extended to a number of tasks such as regression [Vapnik *et al.* '97], principal component analysis [Schölkopf *et al.* '99], etc.
- Most popular optimization algorithms for SVMs are SMO [Platt '99] and SVM^{light} [Joachims' 99], both use *decomposition* to hill-climb over a subset of α_i 's at a time.
- Tuning SVMs remains a black art: selecting a specific kernel and parameters is usually done in a try-and-see manner.

SVM Extensions

- Regression
- Variable Selection
- Boosting
- Density Estimation
- Unsupervised Learning
 - Novelty/Outlier Detection
 - Feature Detection
 - Clustering

Support Vector Machine Resources

- **SVM Application List**
<http://www.clopinet.com/isabelle/Projects/SVM/applist.html>
- **Kernel machines**
<http://www.kernel-machines.org/>
- **Pattern Classification and Machine Learning**
<http://clopinet.com/isabelle/#projects>
- **R a GUI language for statistical computing and graphics**
<http://www.r-project.org/>
- **Kernel Methods for Pattern Analysis – 2004**
<http://www.kernel-methods.net/>
- **An Introduction to Support Vector Machines**
(and other kernel-based learning methods)
<http://www.support-vector.net/>
- **Kristin P. Bennett web page**
<http://www.rpi.edu/~bennek>
- **Isabelle Guyon's home page**
<http://clopinet.com/isabelle>

Support Vector Machine References

- Duda R.O. and Hart P.E.; *Pattern Classification and Scene Analysis*. Wiley, 1973.
- T.M. Cover. Geometrical and statistical properties of systems of linear inequalities with applications in pattern recognition. *IEEE Transactions on Electronic Computers*, 14:326--334, 1965.
- V.Vapnik and A.Lerner. Pattern recognition using generalized portrait method. *Automation and Remote Control*, 24, 1963.
- V.Vapnik and A.Chervonenkis. A note on one class of perceptrons. *Automation and Remote Control*, 25, 1964.
- J.K. Anlauf and M.Biehl. The adatron: an adaptive perceptron algorithm. *Europhysics Letters*, 10:687--692, 1989.
- N.Aronszajn. Theory of reproducing kernels. *Transactions of the American Mathematical Society*, 68:337--404, 1950.
- M.Aizerman, E.Braverman, and L.Rozonoer. Theoretical foundations of the potential function method in pattern recognition learning. *Automation and Remote Control* 25:821--837, 1964.
- O. L. Mangasarian. Linear and nonlinear separation of patterns by linear programming. *Operations Research*, 13:444--452, 1965.
- F. W. Smith. Pattern classifier design by linear programming. *IEEE Transactions on Computers*, C-17:367--372, 1968.
- C.Cortes and V.Vapnik. Support vector networks. *Machine Learning*, 20:273--297, 1995. V.Vapnik. *The Nature of Statistical Learning Theory*. Springer Verlag, 1995.
- V.Vapnik. *Statistical Learning Theory*. Wiley, 1998. A.N. Tikhonov and V.Y. Arsenin. *Solutions of Ill-posed Problems*. W. H. Winston, 1977.
- B.Schoelkopf, C.J.C. Burges, and A.J. Smola, *Advances in kernel methods - support vector learning*, MIT Press, Cambridge, MA, 1999.
- A.J. Smola, P.Bartlett, B.Schoelkopf, and C.Schuermans, *Advances in large margin classifiers*, MIT Press, Cambridge, MA, 1999.